

```
In [1]: # From: https://www.kaggle.com/datasets/naddamuhamed/sleepy-driver-eeg-brai
# Download data as zip: "acquiredDataset.csv"
```

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats
import pandas as pd
import seaborn as sns

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from sklearn.preprocessing import StandardScaler, OneHotEncoder

from scipy.stats import ttest_ind
```

## Data Overview ( acquiredDataset.csv )

- The dataset contains EEG brainwave data collected from people in both awake and asleep states using a NeuroSky MindWave sensor
- Features include attention scores, meditation (calmness) levels, and various brainwave frequencies (delta, theta, alpha, beta, gamma)
- Binary classification: 0 = Awake, 1 = Sleepy

---

## Analysis Steps & Findings:

```
In [3]: data = pd.read_csv('acquiredDataset.csv')
```

```
In [4]: data.head()
```

```
Out[4]:
```

	attention	meditation	delta	theta	lowAlpha	highAlpha	lowBeta	highBeta	lc
0	26	34	960462	277180	26575	27356	26575	13056	
1	29	54	39145	28225	20172	39551	20172	9933	
2	40	48	75410	43144	8601	13564	8601	11663	
3	66	47	16057	41211	2534	34254	2534	27663	
4	81	67	10304	47239	33158	47349	33158	16328	


# 1. Initial Data Exploration

```
In [5]: # @title EEG waveform frequency
columns = ['delta', 'theta', 'lowAlpha', 'highAlpha', 'lowBeta', 'highBeta',
           'lowGamma', 'highGamma']

# Create subplots
fig, axes = plt.subplots(1, len(columns), figsize=(20, 5), sharey=True)

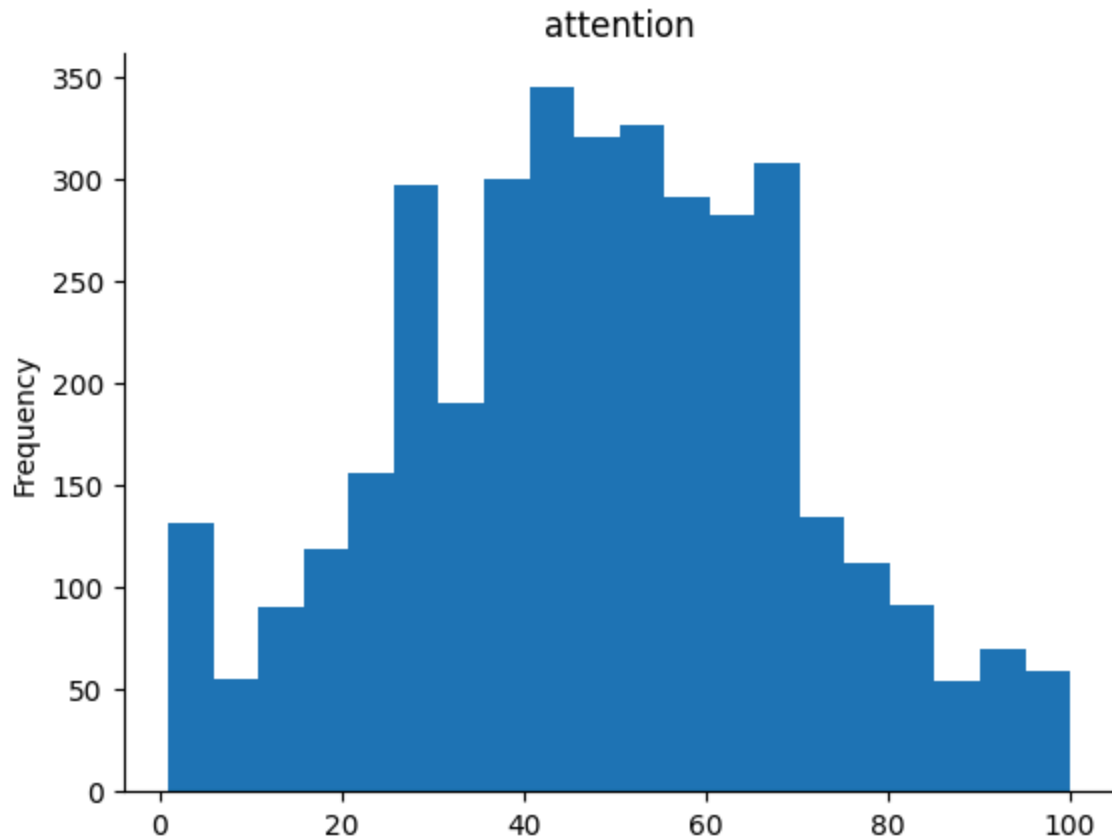
# Plot each column
for ax, col in zip(axes, columns):
    data[col].plot(kind='hist', bins=20, ax=ax, title=col, alpha=0.7)
    ax.spines[['top', 'right']].set_visible(False)
    ax.set_xlabel(col)
    ax.set_ylabel('Frequency')

# Adjust layout
plt.tight_layout()
plt.show()
```



```
In [6]: # @title attention

from matplotlib import pyplot as plt
data['attention'].plot(kind='hist', bins=20, title='attention')
plt.gca().spines[['top', 'right']].set_visible(False)
```



Recap of 1. **Initial Data Exploration**

- Visualized distribution of brainwave frequencies
  - Examined attention scores distribution
- 

## 2. Statistical Analysis

### Are sleep and attention related?

- We will separate the data according to sleepy vs awake - according to our project's hypothesis, the sleepy drivers should have lower attention scores due to poor sleep quality

```
In [7]: group_awake = data[data['classification'] == 0]['attention']
group_sleepy = data[data['classification'] == 1]['attention']

t_stat, p_val = ttest_ind(group_awake, group_sleepy, equal_var=False)

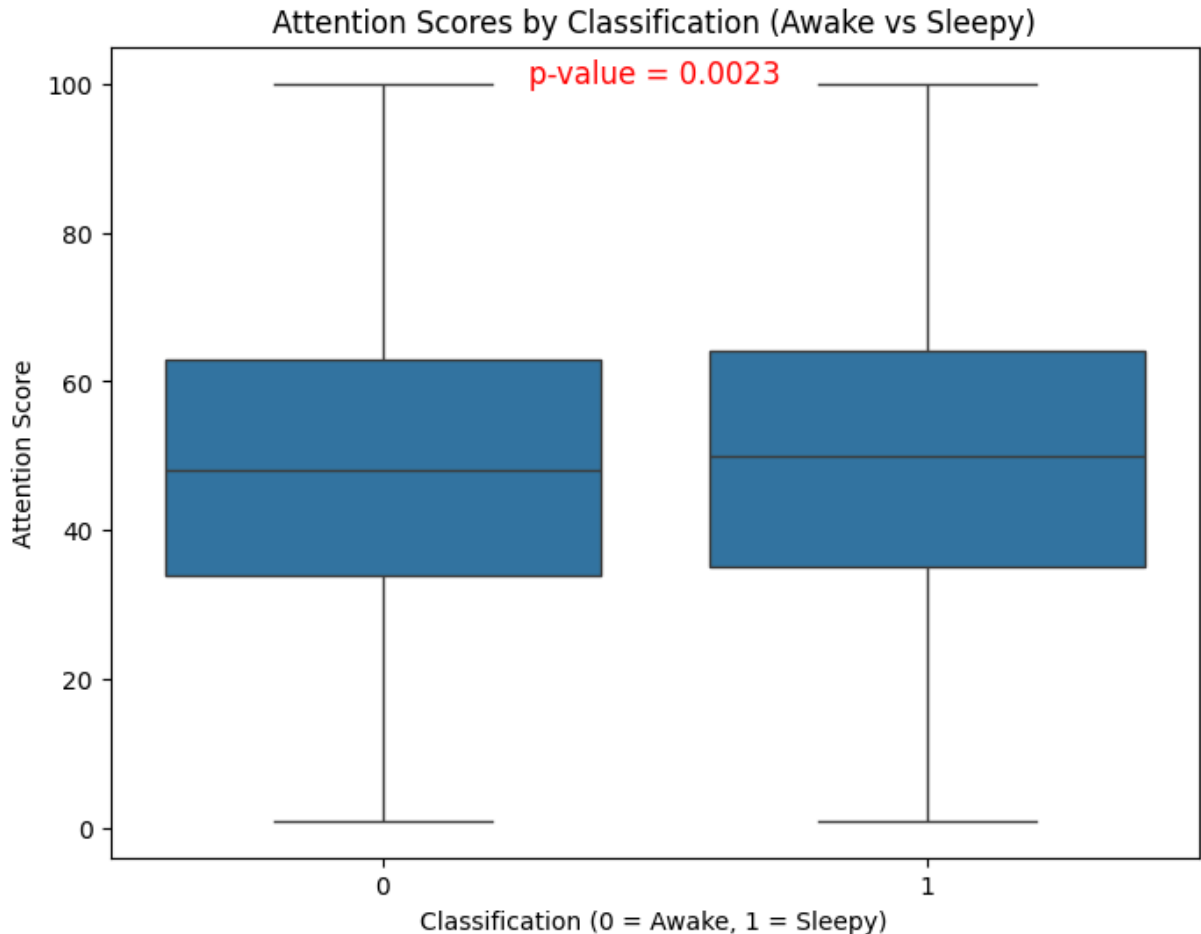
plt.figure(figsize=(8, 6))
sns.boxplot(x='classification', y='attention', data=data)
plt.title("Attention Scores by Classification (Awake vs Sleepy)")

plt.text(0.5, max(data['attention']), f"p-value = {p_val:.4f}",
```

```

horizontalalignment='center', fontsize=12, color='red')
plt.xlabel("Classification (0 = Awake, 1 = Sleepy)")
plt.ylabel("Attention Score")
plt.show()

```



Since  $p=0.0023 < 0.05$ , there is a statistically significant difference between the two groups.

**Now, will the brainwave data differ statistically between the two groups?**

```

In [8]: eeg = ["delta", "theta", "lowAlpha", "highAlpha", "lowBeta", "highBeta", "lc

plt.figure(figsize=(14, 10))
for i, col in enumerate(eeg, 1):
    plt.subplot(4, 2, i)
    sns.boxplot(x='classification', y=col, data=data)
    plt.title(f"{col} by Classification")
    plt.xlabel("Classification (0 = Awake, 1 = Sleepy)")
    plt.ylabel(col)
plt.tight_layout()
plt.show()

results = {}
for col in eeg:
    group_aware = data[data['classification'] == 0][col]

```

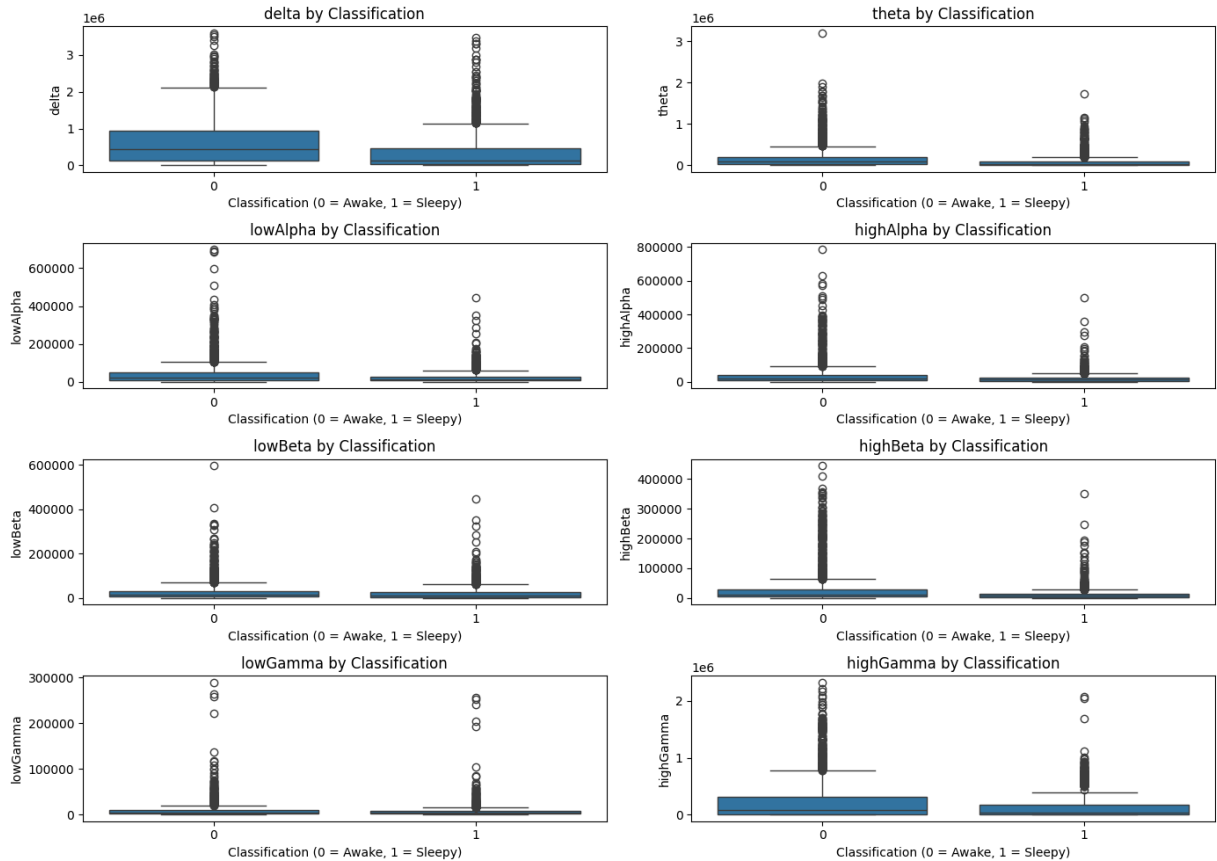
```

group_sleepy = data[data['classification'] == 1][col]

t_stat, p_val = ttest_ind(group_awake, group_sleepy, equal_var=False)
results[col] = {'t_stat': t_stat, 'p_val': p_val}

print("Statistical Test Results:")
for feature, result in results.items():
    print(f"{feature}: t-stat = {result['t_stat']:.4f}, p-value = {result['p_val']:.4f}")

```



Statistical Test Results:

delta: t-stat = 14.8369, p-value = 2.0308e-48

theta: t-stat = 13.8615, p-value = 1.4658e-42

lowAlpha: t-stat = 11.7861, p-value = 1.9027e-31

highAlpha: t-stat = 13.1178, p-value = 2.7364e-38

lowBeta: t-stat = 4.0104, p-value = 6.1816e-05

highBeta: t-stat = 14.9164, p-value = 1.7063e-48

lowGamma: t-stat = 3.0203, p-value = 2.5433e-03

highGamma: t-stat = 6.6961, p-value = 2.4624e-11


Based on statistical analysis, it seems as though **high beta and delta waves** are the most statistically different amongst sleepy vs. awake individuals, closely followed by **delta and theta waves**. This validates our literature review which states "Fast frequencies correspond to beta (13 to 25) and gamma (25 to 60 Hz) waves. They are associated with a high state of vigilance, or cognitive activity. Slow waves correspond to theta (4 to 8 Hz) and delta waves (1 to 4 Hz), and are associated with a state of drowsiness and sleep, respectively." (Terlow 2016).

Recap of 2. **Statistical Analysis**

- Conducted t-tests between awake/sleepy groups
- Found significant differences in attention scores ( $p=0.0023$ )
- High beta and delta waves showed strongest statistical differences between states
- Results aligned with literature: fast frequencies (beta/gamma) indicate vigilance, slow waves (theta/delta) indicate drowsiness

### 3. Feature Relationships

```
In [9]: correlation_matrix = data.corr()
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap="coolwarm", cbar=True)
plt.title("Correlation Heatmap of Features")
plt.show()
```




```
In [10]: selected_columns = [
    "delta", "theta", "lowAlpha", "highAlpha",
    "lowBeta", "highBeta", "lowGamma", "highGamma", "attention"
]

missing_columns = [col for col in selected_columns if col not in data.columns]
if missing_columns:
    print(f"Missing columns in dataset: {missing_columns}")
else:
```

```
sns.pairplot(data[selected_columns], diag_kind="kde", corner=True, plot_
plt.suptitle("Pairwise Relationships Between Brainwave Features and Atte
plt.show()
```

Pairwise Relationships Between Brainwave Features and Attention



### Recap of 3. Feature Relationships

- Created correlation heatmap
- Generated pairwise plots to visualize relationships between brainwaves and attention

## 4. Predictive Modeling

# Can we Predict Attention (Cognitive Performance) based on EEG signals from sleep?

- We will train a few regression models as well as a neural network and compare their performances.
- Since we proved strong statistically significant differences between brainwaves and sleep state, we will train the model using the classification column: it will function as a one hot encoder.

```
In [11]: eeg = data[['delta', 'theta', 'lowAlpha', 'highAlpha', 'lowBeta', 'highBeta']]
target = data['attention']
encoder = OneHotEncoder(drop='first', sparse_output=False) # drop='first' as we don't need the first column
classification_encoded = encoder.fit_transform(data[['classification']])
classification_df = pd.DataFrame(classification_encoded, columns=['classification'])
eeg_encode = pd.concat([eeg, classification_df], axis=1) # Concatenate data
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(eeg_encode, target, test_size=0.2, random_state=42)
```

```
In [13]: # scaler = StandardScaler()
# X_train_scaled = scaler.fit_transform(X_train)
# X_test_scaled = scaler.transform(X_test)
```

```
In [14]: models = {
    "Linear Regression": LinearRegression(),
    "Random Forest Regressor": RandomForestRegressor(random_state=42),
    "Gradient Boosting Regressor": GradientBoostingRegressor(random_state=42),
    "Support Vector Regressor": SVR(kernel='rbf'),
    "K-Nearest Neighbors Regressor": KNeighborsRegressor()
}
```

```
In [15]: results = {}
```

```
In [16]: for model_name, model in models.items():

    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    results[model_name] = {"MSE": mse, "R²": r2, "Predictions": y_pred}
```

```
In [17]: fig, axes = plt.subplots(3, 2, figsize=(14, 12))
axes = axes.flatten()

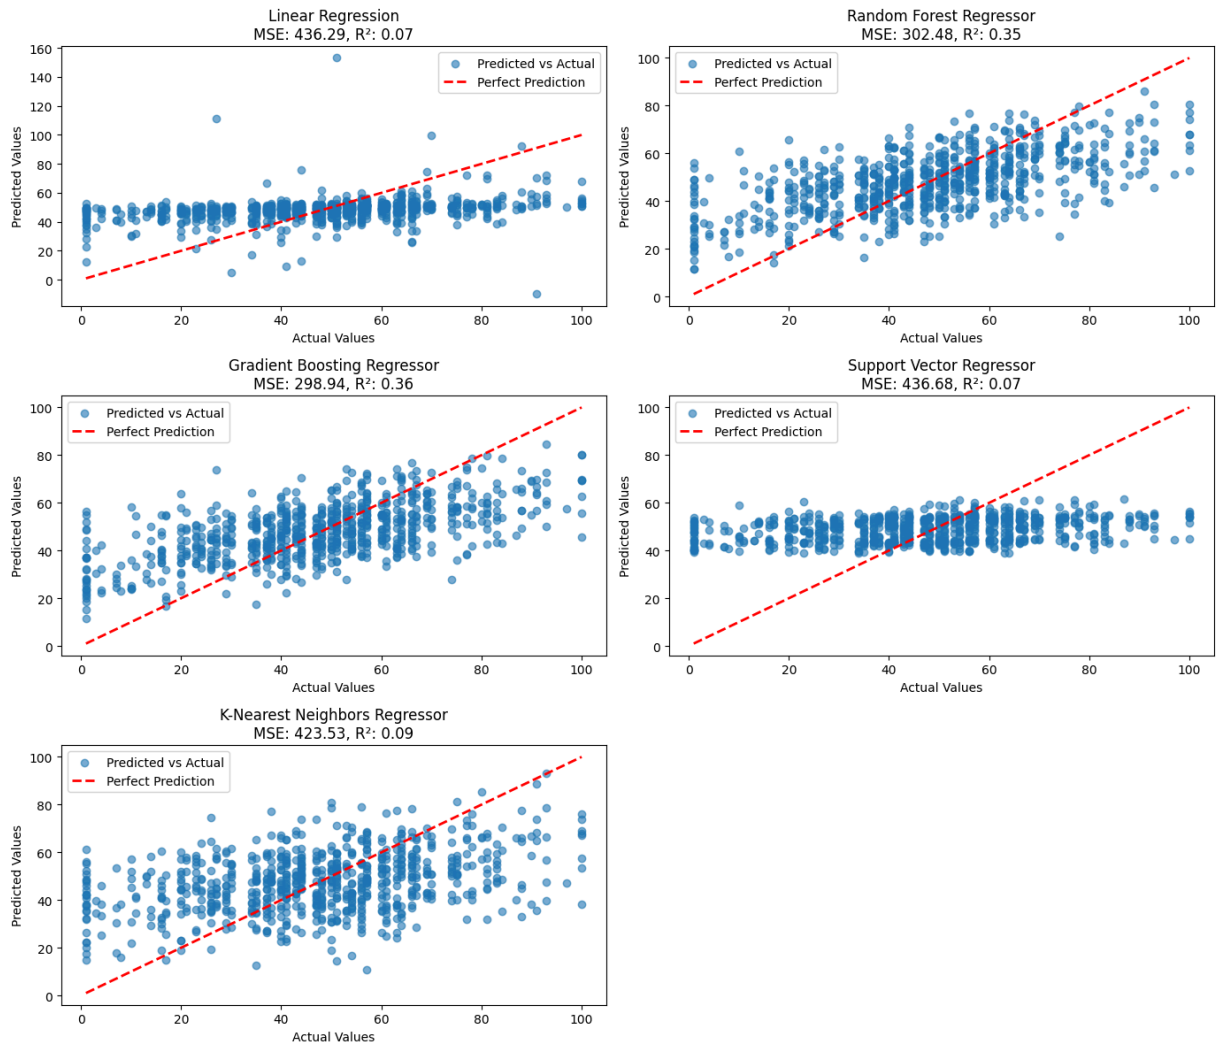
for i, (model_name, result) in enumerate(results.items()):
    ax = axes[i]
    ax.scatter(y_test, result["Predictions"], alpha=0.6, label='Predicted vs Actual')
    ax.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
    ax.set_title(f'{model_name}\nMSE: {result["MSE"]:.2f}, R²: {result["R²"]:.2f}')
    ax.set_xlabel("Actual Values")
    ax.set_ylabel("Predicted Values")
```



```
ax.legend()

plt.tight_layout()
for j in range(len(results), len(axes)):
    axes[j].axis('off')

plt.show()
```



## Model Performance Summary

Model	MSE	R <sup>2</sup>
Linear Regression	436.29	0.07
Random Forest	302.48	0.35
Gradient Boosting	298.94	0.36
SVR	436.68	0.07
KNN	423.53	0.09

Best performing model: Gradient Boosting Regressor

→ The Gradient Boosting Regressor performed the highest with a high  $R^2$  score and lowest MSE.

## Neural Network

```
In [18]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
In [19]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [20]: model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train_scaled.shape[1],)),
    Dropout(0.2),
    Dense(32, activation='relu'),
    Dense(1)
])

model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
model.summary()
```

/Users/raviriley/Library/Caches/pypoetry/virtualenvs/sleep--5W40\_9x-py3.10/lib/python3.10/site-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Par
dense (Dense)	(None, 64)	
dropout (Dropout)	(None, 64)	
dense_1 (Dense)	(None, 32)	2
dense_2 (Dense)	(None, 1)	


Total params: 2,753 (10.75 KB)


Trainable params: 2,753 (10.75 KB)


Non-trainable params: 0 (0.00 B)


```
In [21]: history = model.fit(
    X_train_scaled, y_train,
    validation_split=0.2,
```


```
epochs=50,  
batch_size=32,  
verbose=1  
)
```


Epoch 1/50  
75/75  1s 2ms/step - loss: 2735.3459 - mae: 47.4645 - val\_loss: 2258.6592 - val\_mae: 42.4439


Epoch 2/50  
75/75  0s 922us/step - loss: 1957.2992 - mae: 38.6921 - val\_loss: 886.0980 - val\_mae: 24.3095


Epoch 3/50  
75/75  0s 889us/step - loss: 748.3942 - mae: 22.1919 - val\_loss: 634.3730 - val\_mae: 20.1251


Epoch 4/50  
75/75  0s 876us/step - loss: 628.5749 - mae: 19.8538 - val\_loss: 562.7963 - val\_mae: 18.8466


Epoch 5/50  
75/75  0s 881us/step - loss: 573.1982 - mae: 18.8778 - val\_loss: 515.5413 - val\_mae: 18.0104


Epoch 6/50  
75/75  0s 863us/step - loss: 487.1070 - mae: 17.5834 - val\_loss: 484.9637 - val\_mae: 17.4665


Epoch 7/50  
75/75  0s 893us/step - loss: 497.5269 - mae: 17.6226 - val\_loss: 462.9465 - val\_mae: 17.1019


Epoch 8/50  
75/75  0s 845us/step - loss: 471.6473 - mae: 17.2226 - val\_loss: 448.3519 - val\_mae: 16.8259


Epoch 9/50  
75/75  0s 1ms/step - loss: 442.9675 - mae: 16.7267 - val\_loss: 435.7083 - val\_mae: 16.6388


Epoch 10/50  
75/75  0s 888us/step - loss: 448.9631 - mae: 16.9308 - val\_loss: 426.8712 - val\_mae: 16.5239


Epoch 11/50  
75/75  0s 854us/step - loss: 430.4935 - mae: 16.5378 - val\_loss: 420.9412 - val\_mae: 16.4028


Epoch 12/50  
75/75  0s 919us/step - loss: 427.4755 - mae: 16.6559 - val\_loss: 415.5627 - val\_mae: 16.2951


Epoch 13/50  
75/75  0s 820us/step - loss: 418.5377 - mae: 16.2259 - val\_loss: 411.4618 - val\_mae: 16.2289


Epoch 14/50  
75/75  0s 869us/step - loss: 417.7725 - mae: 16.2187 - val\_loss: 406.9445 - val\_mae: 16.1383



















Epoch 15/50  
75/75  0s 875us/step - loss: 414.7763 - mae: 16.2908 - val\_loss: 405.0710 - val\_mae: 16.0758

Epoch 16/50  
75/75  0s 946us/step - loss: 436.0999 - mae: 16.6309 - val\_loss: 401.6443 - val\_mae: 16.0344

Epoch 17/50  
75/75  0s 890us/step - loss: 429.8110 - mae: 16.6514 - val\_loss: 399.4407 - val\_mae: 15.9647

Epoch 18/50  
75/75  0s 927us/step - loss: 398.1504 - mae: 15.8413 - val\_loss: 397.5006 - val\_mae: 15.9337

Epoch 19/50  
75/75  0s 969us/step - loss: 409.2348 - mae: 16.1791 - val\_loss: 401.6443 - val\_mae: 16.0344

al\_loss: 399.6777 - val\_mae: 16.0000  
Epoch 20/50  
75/75  0s 937us/step - loss: 404.3939 - mae: 16.0681 - v  
al\_loss: 395.2018 - val\_mae: 15.9041  
Epoch 21/50  
75/75  0s 943us/step - loss: 408.3559 - mae: 16.0066 - v  
al\_loss: 393.7819 - val\_mae: 15.8819  
Epoch 22/50  
75/75  0s 915us/step - loss: 400.5138 - mae: 15.9738 - v  
al\_loss: 391.8469 - val\_mae: 15.8006  
Epoch 23/50  
75/75  0s 876us/step - loss: 376.3546 - mae: 15.6291 - v  
al\_loss: 393.3839 - val\_mae: 15.7893  
Epoch 24/50  
75/75  0s 894us/step - loss: 419.7965 - mae: 16.2438 - v  
al\_loss: 389.7506 - val\_mae: 15.7591  
Epoch 25/50  
75/75  0s 952us/step - loss: 402.2542 - mae: 15.9773 - v  
al\_loss: 389.6898 - val\_mae: 15.7816  
Epoch 26/50  
75/75  0s 946us/step - loss: 378.0498 - mae: 15.5387 - v  
al\_loss: 390.7380 - val\_mae: 15.7264  
Epoch 27/50  
75/75  0s 950us/step - loss: 397.7905 - mae: 15.8997 - v  
al\_loss: 387.1385 - val\_mae: 15.6956  
Epoch 28/50  
75/75  0s 883us/step - loss: 403.6649 - mae: 15.9662 - v  
al\_loss: 387.7490 - val\_mae: 15.7545  
Epoch 29/50  
75/75  0s 954us/step - loss: 382.8876 - mae: 15.4865 - v  
al\_loss: 385.8104 - val\_mae: 15.6824  
Epoch 30/50  
75/75  0s 933us/step - loss: 411.3086 - mae: 16.1755 - v  
al\_loss: 383.9679 - val\_mae: 15.6294  
Epoch 31/50  
75/75  0s 922us/step - loss: 397.2262 - mae: 15.8384 - v  
al\_loss: 384.1627 - val\_mae: 15.6072  
Epoch 32/50  
75/75  0s 894us/step - loss: 402.7354 - mae: 16.0752 - v  
al\_loss: 381.7791 - val\_mae: 15.5853  
Epoch 33/50  
75/75  0s 880us/step - loss: 412.0653 - mae: 16.1625 - v  
al\_loss: 380.7258 - val\_mae: 15.5835  
Epoch 34/50  
75/75  0s 897us/step - loss: 420.3224 - mae: 16.3728 - v  
al\_loss: 379.0074 - val\_mae: 15.5341  
Epoch 35/50  
75/75  0s 1ms/step - loss: 402.5019 - mae: 16.0123 - val  
\_loss: 378.3146 - val\_mae: 15.5175  
Epoch 36/50  
75/75  0s 883us/step - loss: 399.3326 - mae: 15.8930 - v  
al\_loss: 378.7667 - val\_mae: 15.5098  
Epoch 37/50  
75/75  0s 933us/step - loss: 381.1002 - mae: 15.6498 - v  
al\_loss: 377.5035 - val\_mae: 15.4959  
Epoch 38/50

```

75/75 ————— 0s 907us/step - loss: 385.8171 - mae: 15.7196 - v
al_loss: 375.6226 - val_mae: 15.4525
Epoch 39/50
75/75 ————— 0s 902us/step - loss: 368.4117 - mae: 15.3309 - v
al_loss: 372.4229 - val_mae: 15.3894
Epoch 40/50
75/75 ————— 0s 898us/step - loss: 383.2478 - mae: 15.6257 - v
al_loss: 372.6351 - val_mae: 15.4003
Epoch 41/50
75/75 ————— 0s 911us/step - loss: 413.9254 - mae: 16.3524 - v
al_loss: 371.8273 - val_mae: 15.3822
Epoch 42/50
75/75 ————— 0s 1ms/step - loss: 372.6580 - mae: 15.2969 - val
_loss: 371.3007 - val_mae: 15.3657
Epoch 43/50
75/75 ————— 0s 975us/step - loss: 391.9673 - mae: 15.8021 - v
al_loss: 371.5344 - val_mae: 15.3765
Epoch 44/50
75/75 ————— 0s 965us/step - loss: 383.0977 - mae: 15.4571 - v
al_loss: 371.2492 - val_mae: 15.3377
Epoch 45/50
75/75 ————— 0s 1ms/step - loss: 371.2884 - mae: 15.3169 - val
_loss: 369.0620 - val_mae: 15.2968
Epoch 46/50
75/75 ————— 0s 905us/step - loss: 381.9666 - mae: 15.5146 - v
al_loss: 365.8521 - val_mae: 15.2631
Epoch 47/50
75/75 ————— 0s 878us/step - loss: 380.6219 - mae: 15.7021 - v
al_loss: 365.5903 - val_mae: 15.2376
Epoch 48/50
75/75 ————— 0s 876us/step - loss: 374.9511 - mae: 15.4957 - v
al_loss: 364.6068 - val_mae: 15.2180
Epoch 49/50
75/75 ————— 0s 911us/step - loss: 376.6045 - mae: 15.4717 - v
al_loss: 364.9396 - val_mae: 15.2102
Epoch 50/50
75/75 ————— 0s 866us/step - loss: 365.2273 - mae: 15.1515 - v
al_loss: 364.3828 - val_mae: 15.1714

```

```

In [22]: test_loss, test_mae = model.evaluate(X_test_scaled, y_test, verbose=0)
         print(f"Test Loss (MSE): {test_loss}, Test MAE: {test_mae}")

```


Test Loss (MSE): 356.1883239746094, Test MAE: 14.772906303405762

```

In [23]: import matplotlib.pyplot as plt

         plt.plot(history.history['loss'], label='Training Loss')
         plt.plot(history.history['val_loss'], label='Validation Loss')
         plt.title('Model Loss')
         plt.xlabel('Epochs')
         plt.ylabel('Loss')
         plt.legend()
         plt.show()


```



```
In [24]: y_pred = model.predict(X_test_scaled)

plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs. Predicted')
plt.show()
```

24/24 ————— 0s 1ms/step



#### Recap of 4. **Predictive Modeling**

- Goal: Predict attention based on EEG signals
  - Tested multiple regression models:
    - Linear Regression
    - Random Forest
    - Gradient Boosting (best performer)
    - SVR
    - KNN
  - Implemented Neural Network:
    - 2-layer architecture with dropout
    - Used StandardScaler for feature normalization
    - Monitored training/validation loss
- 

## Conclusion

### Key Findings

1. Significant difference in attention scores between awake and sleepy states ( $p=0.0023$ )



2. Gradient Boosting achieved best prediction performance

## Limitations

1. Single-channel EEG data
2. Limited sample size
3. Controlled environment

## Future Work

1. Collect multi-channel EEG data
2. Could add feature importance analysis
3. Could do cross-validation for more robust model evaluation
4. Might benefit from hyperparameter tuning
5. Could explore more advanced signal processing techniques for EEG data